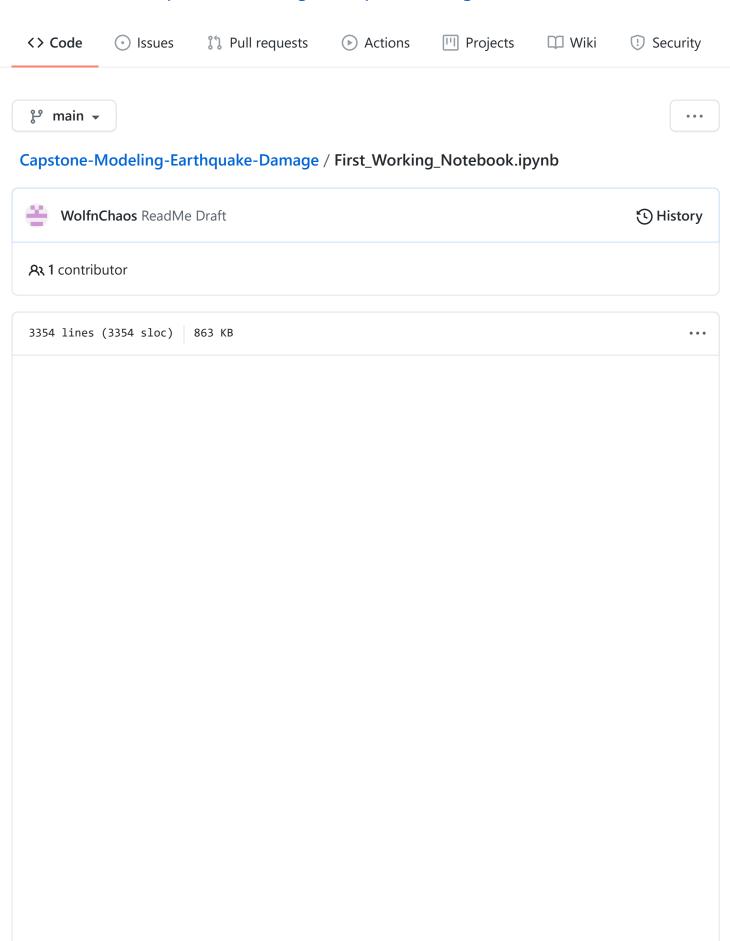
☐ WolfnChaos / Capstone-Modeling-Earthquake-Damage



Overview

A magnitude 7.8 earthquake struck Nepal on April 25, 2015, toppling multi-story buildings in Kathmandu, the capital, and creating landslides and avalanches in the Himalaya Mountains. Nearly 9,000 people died and more than 22,000 suffered injuries.

The quake was followed by hundreds of aftershocks, and only 17 days later, there was another major quake, a magnitude 7.3 temblor. Thirty-nine of the nation's 75 districts with a population of 8 million people — about a third of the national population — were affected. Hundreds of thousands of people lost everything and faced extreme poverty.

More than 600,000 homes were destroyed and more than 288,000 were damaged in the 14 worst-hit districts. Hundreds of thousands of people lost everything and faced extreme poverty and homeless.

Business Problem

The Federal Democratic Republic of Nepal wants to avoid future building damages by reinforcing homes/buildings. They are wanting to know the possible damage risk level that current homes/buildings are at. So, they can better focus their resources, and protect their citizens of Nepal if and when another major earthquake occurs.

Data Understanding

The data that will be used to predict the damage risk level comes from https://www.drivendata.org/), while original data comes from http://eq2015.npc.gov.np/#/). One of the largest dataset done on the aftermath of an earthquake.

The dataset mainly consists of 260601 rows each with information on the building structure and ownership. There are 40 columns in this dataset, where the building_id column is a unique and the target. The remaining 38 features are described in the section below.

Driven Data also obfuscated random lowercase ascii characters to the categorical variables. Using the some of the original data we should be able to find out what this varilbes are and hopefully get some insight when doing the Exploratory Data Analysis (EDA).

Target

Target	Info
Grade 1	represents low damage
Grade 2	represents a medium amount of damage
Grade 3	represents almost complete destruction

Features

Feature	Info
geo_level_1_id, geo_level_2_id, geo_level_3_id (type: int):	geographic region in which building exists, from largest (level 1) to most specific subregion (level 3). Possible values: level 1: 0-30, level 2: 0-1427, level 3: 0-12567.
count_floors_pre_eq (type: int):	number of floors in the building before the earthquake.
age (type: int):	age of the building in years.
area_percentage (type: int):	normalized area of the building footprint.
height_percentage (type: int):	normalized height of the building footprint.
land_surface_condition (type: categorical):	surface condition of the land where the

Zapstorie-wodening-Earthquake-Damage/F	built. Possible values: n, o, t.
foundation_type (type: categorical):	type of foundation used while building. Possible values: h, i, r, u, w.
roof_type (type: categorical):	type of roof used while building. Possible values: n, q, x.
ground_floor_type (type: categorical):	type of the ground floor. Possible values: f, m, v, x, z.
other_floor_type (type: categorical):	type of constructions used in higher than the ground floors (except of roof). Possible values: j, q, s, x.
position (type: categorical):	position of the building. Possible values: j, o, s, t.
plan_configuration (type: categorical):	building plan configuration. Possible values: a, c, d, f, m, n, o, q, s, u.
has_superstructure_adobe_mud (type: binary):	flag variable that indicates if the superstructure was made of

	Adobe/Mud.
has_superstructure_mud_mortar_stone (type: binary):	flag variable that indicates if the superstructure was made of Mud Mortar - Stone.
has_superstructure_stone_flag (type: binary):	flag variable that indicates if the superstructure was made of Stone.
has_superstructure_cement_mortar_stone (type: binary):	flag variable that indicates if the superstructure was made of Cement Mortar - Stone.
has_superstructure_mud_mortar_brick (type: binary):	flag variable that indicates if the superstructure was made of Mud Mortar - Brick.
has_superstructure_cement_mortar_brick (type: binary):	flag variable that indicates if the superstructure was made of Cement Mortar - Brick.
has_superstructure_timber (type: binary):	flag variable that indicates if the superstructure was made of Timber.
has_superstructure_bamboo (type: binary):	flag variable that indicates if the superstructure

dapsone medeling Larriquane Burnage.	was made of Bamboo.
has_superstructure_rc_non_engineered (type: binary):	flag variable that indicates if the superstructure was made of non- engineered reinforced concrete.
has_superstructure_rc_engineered (type: binary):	flag variable that indicates if the superstructure was made of engineered reinforced concrete.
has_superstructure_other (type: binary):	flag variable that indicates if the superstructure was made of any other material.
legal_ownership_status (type: categorical):	legal ownership status of the land where building was built. Possible values: a, r, v, w.
count_families (type: int):	number of families that live in the building.
has_secondary_use (type: binary):	flag variable that indicates if the building was used for any secondary purpose.
	flag variable

has_secondary_use_agriculture (type: binary):	that indicates if the building was used for agricultural purposes.	
has_secondary_use_hotel (type: binary):	flag variable that indicates if the building was used as a hotel.	
has_secondary_use_rental (type: binary):	flag variable that indicates if the building was used for rental purposes.	
has_secondary_use_institution (type: binary):	flag variable that indicates if the building was used as a location of any institution.	
has_secondary_use_school (type: binary):	flag variable that indicates if the building was used as a school.	
has_secondary_use_industry (type: binary):	flag variable that indicates if the building was used for industrial purposes.	
has_secondary_use_health_post (type: binary):	flag variable that indicates if the building was used as a health post.	
has_secondary_use_gov_office (type: binary):	flag variable that indicates if the building was used fas a government office.	
		flag

''-	oupsione Modeling Editifiquate Damage/i	ist_vvoikiiig_ivoteboo	ok.ipyrib at ma	
			variable	
			that	
		binary	indicates	
	has accordant use use police		if the	
	has_secondary_use_use_police		building	
			was	
			used as	
			a police	
			station.	
		flag variable		
		that indicates		
		if the building		
	has_secondary_use_other (type: binary):	was		
		secondarily		
		used for other		
		purposes.		

Data Preparation

Import Libraries and Tools

```
In [51]:
         # Import pandas and set column display to m
         ax.
         import pandas as pd
         pd.set option('display.max columns', None)
         # Import matplotlib and seaborn, set style
          theme to whitegrid.
         import matplotlib.pyplot as plt
         import seaborn as sns
         sns.set theme(style="whitegrid")
         # Import pickle
         import pickle
         # Import model and tranformers from sklear
         from sklearn.tree import DecisionTreeClassi
         from sklearn.model_selection import train_t
         est split, GridSearchCV
         from sklearn.metrics import accuracy_score,
         recall score, precision score, f1 score, pl
         ot confusion matrix
         from sklearn.preprocessing import OneHotEnc
         oder, StandardScaler
         from sklearn.compose import ColumnTransform
         er
         from sklearn.impute import SimpleImputer
         from sklearn.neighbors import KNeighborsCla
         ssifier
```

```
from sklearn.naive_bayes import GaussianNB
from sklearn.linear_model import LogisticRe
gression
```

from sklearn.ensemble import RandomForestCl
assifier

from sklearn.inspection import permutation_
importance

Import xgboost classifer.
from xgboost import XGBClassifier

Import SMOTE to handle class imbalacne an d pipeline.

from imblearn.over_sampling import SMOTE
from imblearn.pipeline import Pipeline as P
ipeline

Read in Datasets for Model

Out[2]:

	building_id	geo_level_1_id	geo_level_2_id	ged
0	802906	6	487	121
1	28830	8	900	281
2	94947	21	363	897
3	590882	22	418	106
4	201944	11	131	148
4				•

Checking for Missing Values

In [3]: # Check of missing values.
df.isna().sum()

Out[3]: building_id 0
geo_level_1_id 0
geo_level_2_id 0
geo_level_3_id 0

psione-wodeling-Latinquake-Damage/i iist_working_wotebook.ip	yiib at
count_floors_pre_eq	0
age	0
area_percentage	0
height_percentage	0
<pre>land_surface_condition</pre>	0
foundation_type	0
roof_type	0
<pre>ground_floor_type</pre>	0
other_floor_type	0
position	0
plan_configuration	0
has_superstructure_adobe_mud	0
has_superstructure_mud_mortar_stone	0
has_superstructure_stone_flag	0
has_superstructure_cement_mortar_stone	0
has_superstructure_mud_mortar_brick	0
has_superstructure_cement_mortar_brick	0
has_superstructure_timber	0
has_superstructure_bamboo	0
has_superstructure_rc_non_engineered	0
has_superstructure_rc_engineered	0
has_superstructure_other	0
legal_ownership_status	0
count_families	0
has_secondary_use	0
has_secondary_use_agriculture	0
has_secondary_use_hotel	0
has_secondary_use_rental	0
has_secondary_use_institution	0
has_secondary_use_school	0
has_secondary_use_industry	0
has_secondary_use_health_post	0
has_secondary_use_gov_office	0
has_secondary_use_use_police	0
has_secondary_use_other	0
damage_grade	0
dtype: int64	

No missing values to worry about.

Checking Featuer Types

```
ZOMOMT HOH-HATT THEOA
    geo level 2 id
260601 non-null int64
    geo level 3 id
260601 non-null int64
    count_floors_pre_eq
260601 non-null int64
    age
260601 non-null int64
    area percentage
260601 non-null int64
    height_percentage
260601 non-null int64
    land surface condition
260601 non-null object
    foundation type
260601 non-null object
10 roof_type
260601 non-null object
11 ground floor type
260601 non-null object
12 other_floor_type
260601 non-null object
13 position
260601 non-null object
14 plan configuration
260601 non-null object
15 has_superstructure_adobe_mud
260601 non-null int64
16 has_superstructure_mud_mortar_stone
260601 non-null int64
17 has superstructure stone flag
260601 non-null int64
18 has_superstructure_cement_mortar_stone
260601 non-null int64
19 has_superstructure_mud_mortar_brick
260601 non-null int64
20 has_superstructure_cement_mortar_brick
260601 non-null int64
21 has_superstructure_timber
260601 non-null int64
22 has superstructure bamboo
260601 non-null int64
23 has_superstructure_rc_non_engineered
260601 non-null int64
24 has_superstructure_rc_engineered
260601 non-null int64
25 has_superstructure_other
260601 non-null int64
26 legal_ownership_status
260601 non-null object
27 count families
260601 non-null int64
28 has_secondary_use
260601 non-null int64
29 has secondary use agriculture
260601 non-null int64
30
    has_secondary_use_hotel
```

```
260601 non-null int64
 31 has secondary use rental
260601 non-null int64
 32 has secondary use institution
260601 non-null int64
 33 has secondary use school
260601 non-null int64
 34 has secondary use industry
260601 non-null int64
 35 has_secondary_use_health_post
260601 non-null int64
 36 has_secondary_use_gov_office
260601 non-null int64
 37 has secondary use use police
260601 non-null int64
 38 has_secondary_use_other
260601 non-null int64
 39 damage grade
260601 non-null
                int64
dtypes: int64(32), object(8)
memory usage: 81.5+ MB
```

Some of the features that are mean to be Booleans are classified as int64.

```
In [5]:
        # Making list of all colums that should be
         a boolean type.
        bool list = ['has superstructure adobe mud'
        , 'has_superstructure_mud_mortar_stone', 'h
        as_superstructure_stone_flag',
                      'has superstructure cement mor
        tar stone', 'has superstructure mud mortar
        brick',
                      'has superstructure cement mor
        tar_brick', 'has_superstructure_timber', 'h
        as_superstructure_bamboo',
                      'has superstructure rc non eng
        ineered', 'has superstructure rc engineere
        d', 'has_superstructure_other',
                      'has secondary use', 'has seco
        ndary_use_agriculture', 'has_secondary_use_
        hotel', 'has_secondary_use_rental',
                      'has secondary use institutio
        n', 'has_secondary_use_school', 'has_second
        ary use industry',
                      'has_secondary_use_health_pos
        t', 'has secondary use gov office', 'has se
        condary_use_use_police',
                      'has_secondary_use_other']
        # Loop though list to change types to boole
        an.
        for name in bool list:
            df[name] = df[name].astype('bool')
        # Checking types one more time.
        df.info()
```

w. /

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 260601 entries, 0 to 260600
Data columns (total 40 columns):
     Column
Non-Null Count
                Dtype
    -----
    building id
260601 non-null int64
     geo_level_1_id
260601 non-null int64
    geo level 2 id
260601 non-null int64
     geo_level_3_id
260601 non-null int64
    count floors pre eq
260601 non-null int64
     age
260601 non-null int64
     area_percentage
260601 non-null int64
    height percentage
260601 non-null int64
     land surface condition
260601 non-null object
    foundation_type
260601 non-null object
 10 roof type
260601 non-null object
 11 ground floor type
260601 non-null object
 12 other_floor_type
260601 non-null object
 13 position
260601 non-null object
 14 plan_configuration
260601 non-null object
 15 has_superstructure_adobe_mud
260601 non-null bool
 16 has superstructure mud mortar stone
260601 non-null bool
 17 has_superstructure_stone_flag
260601 non-null bool
 18 has_superstructure_cement_mortar_stone
260601 non-null bool
 19 has superstructure mud mortar brick
260601 non-null bool
 20 has_superstructure_cement_mortar_brick
260601 non-null bool
 21 has_superstructure_timber
260601 non-null bool
 22 has superstructure bamboo
260601 non-null bool
 23 has_superstructure_rc_non_engineered
260601 non-null bool
 24 has_superstructure_rc_engineered
260601 non-null hool
```

```
ZUUUUT HUH HUTT
 25 has superstructure other
260601 non-null bool
 26 legal_ownership_status
260601 non-null object
 27 count families
260601 non-null
                int64
 28 has secondary use
260601 non-null bool
 29 has_secondary_use_agriculture
260601 non-null bool
 30 has secondary use hotel
260601 non-null bool
 31 has secondary use rental
260601 non-null bool
 32 has_secondary_use_institution
260601 non-null bool
 33 has secondary use school
260601 non-null bool
 34 has secondary use industry
260601 non-null bool
 35 has secondary use health post
260601 non-null bool
 36 has secondary use gov office
260601 non-null bool
 37 has_secondary_use_use_police
260601 non-null bool
 38 has secondary use other
260601 non-null bool
 39 damage grade
260601 non-null int64
dtypes: bool(22), int64(10), object(8)
memory usage: 43.2+ MB
```

Categorical Column Value Names Data

All the categorical columns are represented by an obfuscated random lowercase ascii character, will use the some of the original data set to find out what the letters represent.

```
In [6]: # Read in building structure data.
    df_building_structure = pd.read_csv('Data/c
    sv_building_structure.csv')

# Read in building owner one use data set.
    df_building_ownership_and_use = pd.read_csv
    ('Data/csv_building_ownership_and_use.csv')

# Merge the structure and owner into one da
    ta set.
    df_og = df_building_structure.merge(df_buil
    ding_ownership_and_use, on=['building_id',
    'district_id',

    'vdcmun_id', 'ward_id'])

# Or bood()
```

Out[6]:

	building_id	district_id	vdcmun_id	ward_id
0	120101000011	12	1207	120703
1	120101000021	12	1207	120703
2	120101000031	12	1207	120703
3	120101000041	12	1207	120703
4	120101000051	12	1207	120703
4				

In [7]: # Making a list of all categorical columns. categorical_list = ['land_surface_conditio n', 'foundation_type', 'roof_type', 'ground _floor_type', 'other_floor_type', 'po sition', 'plan_configuration', 'legal_owner ship status'] # Loop though each name in list and print t he counts for each value. for name in categorical_list: print('----') print(name) print('***') print(df_og[name].value_counts()) print('***') print(df[name].value_counts())

```
land_surface_condition
Flat
                  631675
Moderate slope
                  105640
Steep slope
                    24791
Name: land surface condition, dtype: int64
     216757
t
      35528
n
       8316
Name: land_surface_condition, dtype: int64
foundation_type
Mud mortar-Stone/Brick
                           628716
Bamboo/Timber
                            57473
Cement-Stone/Brick
                            39245
RC
                            32120
0ther
                             4552
```

```
Name: foundation type, dtype: int64
***
r
     219196
      15118
W
      14260
u
i
      10579
       1448
h
Name: foundation type, dtype: int64
roof_type
***
Bamboo/Timber-Light roof
                             503748
Bamboo/Timber-Heavy roof
                             213774
                              44584
RCC/RB/RBC
Name: roof_type, dtype: int64
***
     182842
n
      61576
q
      16183
Х
Name: roof type, dtype: int64
ground_floor_type
***
               618217
Mud
RC
                73149
Brick/Stone
                 66093
Timber
                  3594
Other
                 1053
Name: ground_floor_type, dtype: int64
***
f
     209619
      24877
Х
V
      24593
       1004
Z
m
Name: ground_floor_type, dtype: int64
other_floor_type
TImber/Bamboo-Mud
                      486907
Timber-Planck
                      123635
Not applicable
                      118822
RCC/RB/RBC
                       32742
Name: other_floor_type, dtype: int64
***
     165282
q
      43448
Х
j
      39843
      12028
Name: other_floor_type, dtype: int64
position
***
Not attached
                    604453
Attached-1 side
                    129432
Attached-2 side
                     26910
Attached-3 side
                     1310
Name: position, dtvpe: int64
```

```
***
     202090
s
t
      42896
j
      13282
       2333
Name: position, dtype: int64
plan configuration
Rectangular
                                      731257
Square
                                       17576
                                       10079
L-shape
T-shape
                                         969
                                         940
Multi-projected
Others
                                         518
U-shape
                                         448
E-shape
                                         140
Building with Central Courtyard
                                          98
                                          80
H-shape
Name: plan_configuration, dtype: int64
***
     250072
d
       5692
q
       3649
u
s
        346
        325
c
        252
а
        159
0
         46
m
         38
n
         22
Name: plan_configuration, dtype: int64
legal_ownership_status
Private
                  731387
Public
                   19232
Institutional
                    7823
0ther
                    3664
Name: legal_ownership_status, dtype: int64
     250939
٧
       5512
а
       2677
       1473
```

Look at the print out it looks like most of the random values in the df dataset do correspond to the df_og dataset. Will use the values to replace the random values.

Name: legal ownership status, dtype: int64

```
e and loop thought the value count index fr
om the df dataset and replace the
    the value with the value count index fo
r the df_og data set.
    '''
    for i in range(len(df[name].value_count
s().index)):
        df[name].replace({df[name].value_co
unts().index[i]:

df_og[name].value_counts().index[i]}, inpla
ce=True)

# Look though each value in list and replac
e them.
for cat_value in categorical_list:
    replace_categorical_value(cat_value)

df.head()
```

Out[8]:

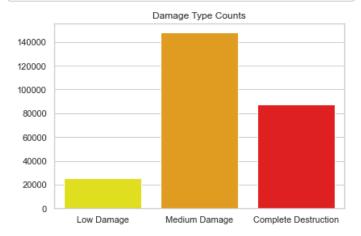
		building_id	geo_level_1_id	geo_level_2_id	geo
	0	802906	6	487	121
	1	28830	8	900	281
	2	94947	21	363	897
	3	590882	22	418	106
-	4	201944	11	131	148
L	◀				•

Exploratory Data Analysis (EDA)

Class Counts

Lets take a look at are target classes and see what kind of counts we have.

```
# making a plot to visualize ladel counts.
יות [דת].
         fig, ax = plt.subplots()
         # Defining x \& y values.
         x = df['damage_grade'].value_counts().index
         y = df['damage_grade'].value_counts().value
         # Making list for label names and colors.
         labels = ['Low Damage', 'Medium Damage', 'C
         omplete Destruction']
         color = ['yellow','orange', 'red']
         #Making bar plot.
         ax = sns.barplot(x=x, y=y, palette=color)
         ax.set xticklabels(labels)
         ax.set_title('Damage Type Counts')
         ax.set_facecolor('white')
         plt.tight_layout()
         #Saving plot to Images folder.
         fig.savefig('Images/classes.png')
```



Looks like are classes are imbalance this may cuase an issues down the road when we start training are models.

Ploting Categorical Columns

```
# Making figure.
fig, ax = plt.subplots()

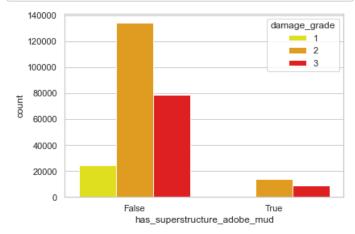
# List of colors.
color = ['yellow','orange', 'red']

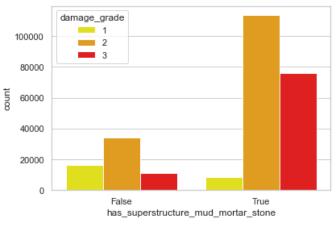
# Making countplot.
sns.countplot(x=name, hue='damage_g
rade', palette=color, data=df, ax=ax)
plt.xticks(rotation=rotation)
plt.tight_layout()

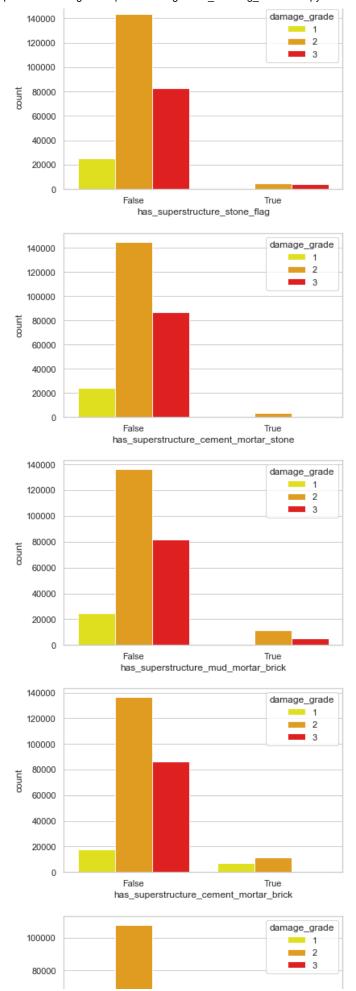
#Saving plot ot Image folder.
fig.savefig(f'Images/{name}.png')
```

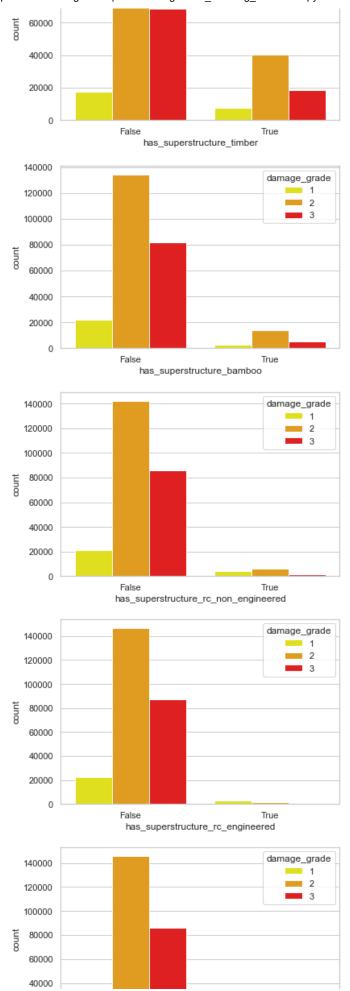
In []:

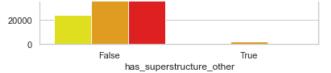
Superstructure Columns











Secondary Columns

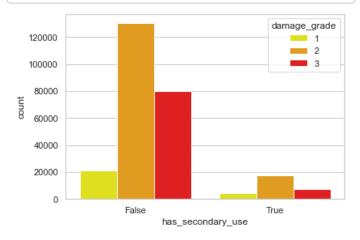
In [162]:

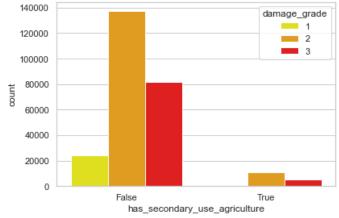
Making list of only the secondary use col umns.

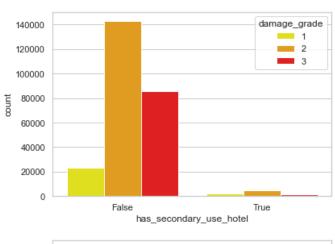
column_secondary = list(df.columns)[28:-1]

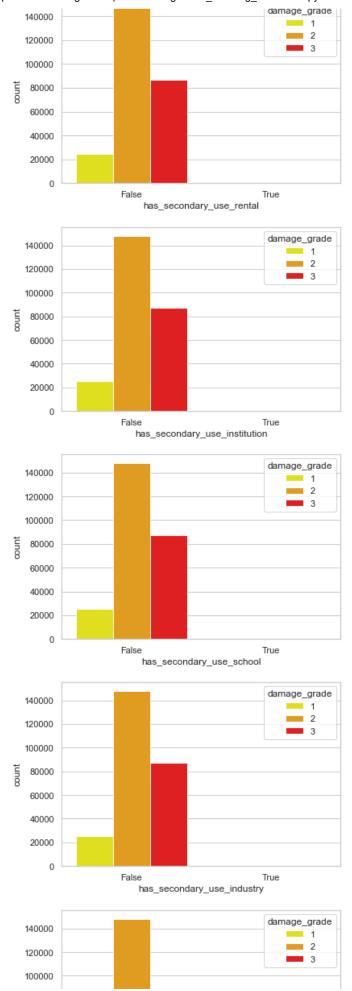
Using function to visualize columns as a countplot.

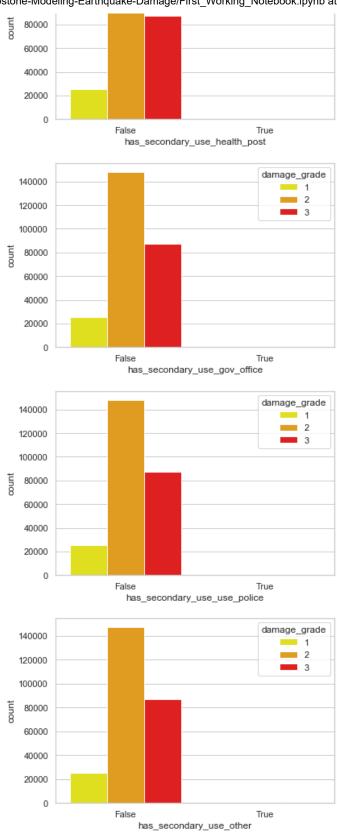
countplot_loop(column_secondary, 0)





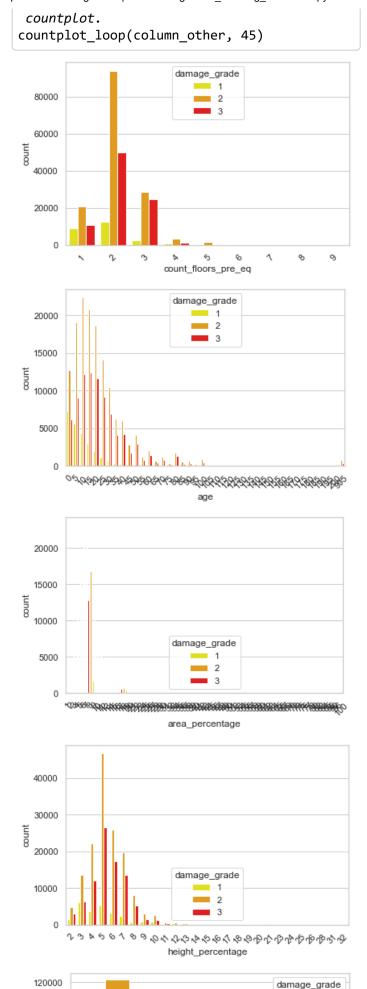


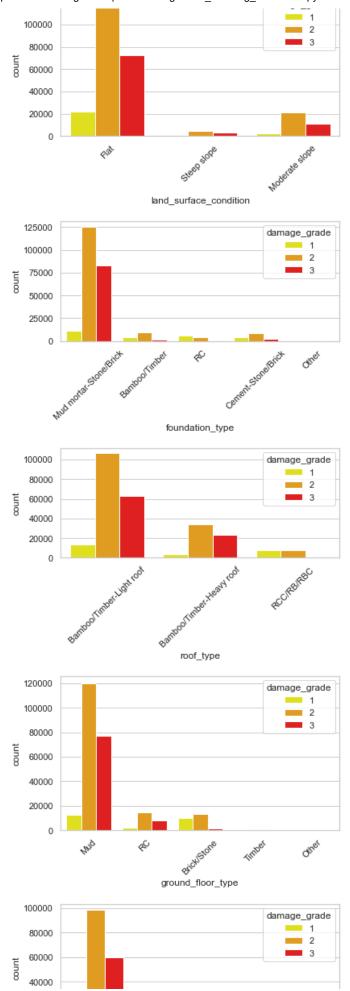


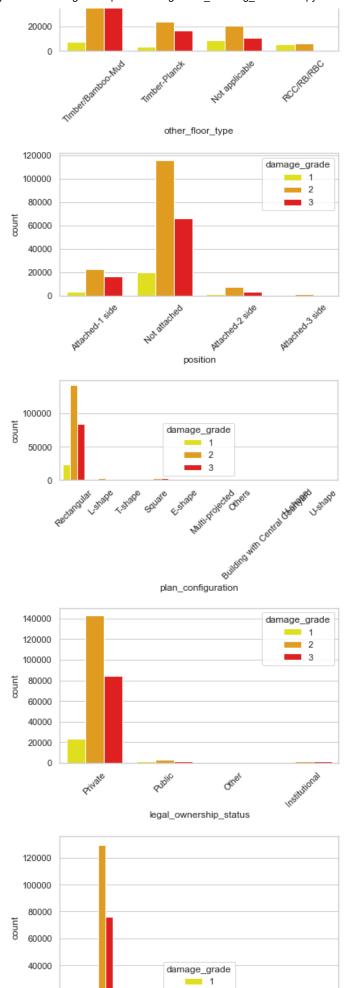


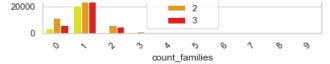
Other Columns

```
In [163]: # Making list of other columns.
    column_other = list(df.columns)[4:15] + lis
    t(df.columns)[26:28]
# Using function to visualize columns as a
```









First Simple Model (FSM)

For the FSM we will use just the numerical and boolean columns and the classifier we will use a Decision Tree Classifer. First things first, we need to split that data by doing a train test split.

Difine the X and y values for fsm.

```
X fsm = df.drop(['building id', 'land surfa
         ce_condition', 'foundation_type',
                   'roof_type', 'ground_floor_type',
         'other_floor_type', 'position', 'plan_confi
         guration',
                   'legal ownership status', 'damage
         grade'], axis=1)
         y_fsm = df['damage_grade']
         # Preform the train test split for fsm.
         X_train_fsm, X_test_fsm, y_train_fsm, y_tes
         t fsm = train test split(X fsm, y fsm, rand
         om state=42, stratify=y fsm)
In [16]: # Instantiate decision tree classifier for
          a fsm.
         fsm dt = DecisionTreeClassifier(random stat
         e = 42)
         # Fir the X and y train date to the fsm.
         fsm dt.fit(X train fsm, y train fsm)
Out[16]: DecisionTreeClassifier(random_state=42)
In [19]:
         # Making function for printing scores.
         def print scores(model, X train, X test, y
         train, y_test):
             This function will print the scoure for
         both train and test.
             # Predicting labels for both train and
          test data.
             y_hat_train = model.predict(X_train)
             y_hat_test = model.predict(X_test)
             # Print out scores.
             print(f'
                             Model Scores')
             print('----')
             print(' Accuracy:', accuracy_score(y_tr
         ain, y_hat_train))
             print('
                     Recall:', recall_score(y_trai
```

```
n, y_hat_train, average = 'macro'))
             print('Precision:', precision_score(y_t
         rain, y_hat_train, average = 'macro'))
             print('
                          F1:', f1_score(y_train, y
         _hat_train, average = 'macro'))
             print('-----')
             print(' Accuracy:', accuracy_score(y_te
         st, y hat test))
             print('
                       Recall:', recall_score(y_test
         , y_hat_test, average = 'macro'))
             print('Precision:', precision_score(y_t
         est, y hat test, average = 'macro'))
             print('
                           F1:', f1_score(y_test, y_
         hat_test, average = 'macro'))
In [20]:
         # Print out fsm scores.
         print_scores(fsm_dt, X_train_fsm, X_test_fs
         m, y train fsm, y test fsm)
                 Model Scores
          Accuracy: 0.9742082374008698
            Recall: 0.9746926003984552
         Precision: 0.969726887943331
                F1: 0.9718795910888217
          Accuracy: 0.6604503384445365
            Recall: 0.6149900294776912
         Precision: 0.6073121797582424
                F1: 0.6109604665814518
```

In [130]:

```
# Making a function of confusion matrix.
def plot_cm(model, X, y):
```

This function with change the sns theme back to plain white,

plot a confusion matrix and then return the sns theme back to whitegrid.

Setting theme back to white for confu sion matrix.

sns.set theme(style="dark")

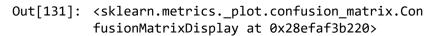
Setting confusion matrix as a variabl e.

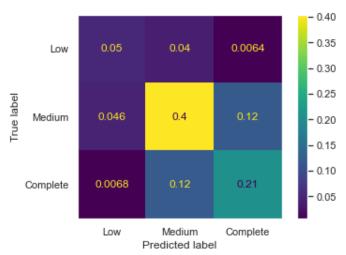
plot = plot_confusion_matrix(model, X, y, normalize='all', display_labels=['Low', 'Medium', 'Complete']);

Returning theme back to whitegrid. sns.set theme(style="whitegrid")

#Return plot. return plot

Plotting fsm confusion matrix. In [131]: plot_cm(fsm_dt, X_test_fsm, y_test_fsm)





For the FSM 16.6% of the testing data is classified as a false negative.

Preprocessing

Setting up preprocessing for vanilla models.

```
In [21]:
         # Difine the X and v values for vms
         X = df.drop(['building id', 'damage grade'
         ], axis=1)
         y = df['damage grade']
         # Preform the train test split for vms.
         X train, X test, y train, y test = train te
         st split(X, y, random state=42, stratify=y)
         # Making a list of all numeric columns.
         num list = ['geo level 1 id', 'geo level 2
         id', 'geo_level_3_id', 'count_floors_pre_e
         q', 'age', 'area percentage',
                      'height percentage', 'count fam
         ilies']
         # Making a list of all boolean columns.
         bool list =['has superstructure adobe mud',
         'has_superstructure_mud_mortar_stone', 'ha
         s superstructure stone flag',
                      'has superstructure cement mort
         ar_stone', 'has_superstructure_mud_mortar_b
         rick',
                      'has superstructure cement mort
         ar brick', 'has superstructure timber', 'ha
         s superstructure bamboo',
                      'has superstructure rc non engi
                   'has_superstructure_rc_engineered'
         neered',
```

```
, nas_superstructure_otner ,
            'has_secondary_use', 'has_secon
dary use agriculture', 'has secondary use h
otel',
            'has_secondary_use_rental', 'ha
s secondary use institution', 'has secondar
y use school',
            'has secondary use industry',
'has secondary use health post', 'has secon
dary_use_gov_office',
            'has_secondary_use_use_police',
'has secondary use other']
# Making a list of all categorical columns.
cat_list = ['land_surface_condition', 'foun
dation_type', 'roof_type', 'ground_floor_ty
pe', 'other_floor_type',
             'position', 'plan_configuratio
n', 'legal ownership status']
# Setting of column transformer for each co
lumn list.
ct = ColumnTransformer([('ohe', OneHotEncod
er(drop='first'), cat_list),
                        ('ss', SimpleImpute
r(), num_list),
                        ('si', SimpleImpute
r(), bool list)])
# Fit and transform the train data and only
transform the test data.
X train ct = ct.fit transform(X train)
X test ct = ct.transform(X test)
# Use smote to handle label imbalance.
sm = SMOTE(random state=42)
X train sm, y train sm = sm.fit sample(X tr
ain_ct, y_train)
```

Vanilla Models

Now we will use the prepocessed data to train and test some vms. The models that we will be using are: KNeighborsClassifier, GaussianNB, LogisticRegression, Random Forest Classifer, and XGB Classifer.

KNeighborsClassifier

```
In [22]: # # Instantiate KNeighborsClassifier and fi
    t X and y.
# knc = KNeighborsClassifier()
# knc.fit(X_train_sm, y_train_sm)
```

```
# # Pickle model to count done when need to
re run notebook.
# with open('Pickle/kneighbors_vm', 'wb') a
s f:
# pickle.dump(knc, f)

# Open pickled model.
with open('Pickle/kneighbors_vm', 'rb') as
f:
    knc = pickle.load(f)

# Print model's scores.
print_scores(knc, X_train_ct, X_test_ct, y_
train_v test)
```